

```
In [2]: import tensorflow as tf
dataset_dir = "C:\Plant_leave_diseases_dataset_without_augmentation"

batch_size = 32
img_size = (224, 224)

dataset = tf.keras.utils.image_dataset_from_directory(
    dataset_dir,
    labels='inferred',
    label_mode='categorical',
    color_mode='rgb',
    batch_size=batch_size,
    image_size=img_size,
    shuffle=True
)

for images, labels in dataset.take(1):
    print("Image batch shape:", images.shape)
    print("Label batch shape:", labels.shape)
```

Found 24164 files belonging to 17 classes.

Image batch shape: (32, 224, 224, 3)

Label batch shape: (32, 17)

```
In [3]: class_names = dataset.class_names
print("Class names:", class_names)
```

Class names: ['Corn__Cercospora_leaf_spot Gray_leaf_spot', 'Corn__Common_rust', 'Corn__Northern_Leaf_Blight', 'Corn__healthy', 'Potato__Early_blight', 'Potato__Late_blight', 'Potato__healthy', 'Tomato__Bacterial_spot', 'Tomato__Early_blight', 'Tomato__Late_blight', 'Tomato__Leaf_Mold', 'Tomato__Septoria_leaf_spot', 'Tomato__Spider_mites Two-spotted_spider_mite', 'Tomato__Target_Spot', 'Tomato__Tomato_Yellow_Leaf_Curl_Virus', 'Tomato__Tomato_mosaic_virus', 'Tomato__healthy']

```
In [4]: dataset_size = tf.data.experimental.cardinality(dataset).numpy()

train_size = int(0.8 * dataset_size)
val_size = int(0.1 * dataset_size)
test_size = dataset_size - train_size - val_size

train_dataset = dataset.take(train_size)
remaining_dataset = dataset.skip(train_size)

val_dataset = remaining_dataset.take(val_size)
test_dataset = remaining_dataset.skip(val_size)

train_dataset = train_dataset.prefetch(buffer_size=tf.data.AUTOTUNE)
val_dataset = val_dataset.prefetch(buffer_size=tf.data.AUTOTUNE)
test_dataset = test_dataset.prefetch(buffer_size=tf.data.AUTOTUNE)

print(f"Total dataset size: {dataset_size}")
print(f"Train dataset size: {train_size}")
print(f"Validation dataset size: {val_size}")
print(f"Test dataset size: {test_size}")
```

Total dataset size: 756
Train dataset size: 604
Validation dataset size: 75
Test dataset size: 77

class balancing

```
In [5]: import matplotlib.pyplot as plt
import numpy as np
from collections import Counter

def get_class_distribution(dataset):
    labels = []
    for _, label in dataset:
        label = label.numpy()
        if label.ndim > 0:
            label = label.argmax()
            labels.append(label)


    class_counts = Counter(labels)
    return class_counts

class_counts = get_class_distribution(train_dataset)

print("Class Distribution:", class_counts)
```

Class Distribution: Counter({14: 125, 7: 64, 9: 48, 13: 44, 16: 43, 12: 40, 11: 34, 3: 34, 5: 33, 10: 30, 1: 25, 2: 23, 8: 19, 4: 18, 0: 12, 15: 10, 6: 2})

```
In [6]: def plot_class_distribution(class_counts):  
    classes = list(class_counts.keys())  
    counts = list(class_counts.values())  
  
    plt.figure(figsize=(8, 5))  
    plt.bar(classes, counts, color='skyblue', alpha=0.7)  
    plt.xlabel("Class Labels")  
    plt.ylabel("Number of Samples")  
    plt.title("Class Distribution in Dataset")  
    plt.xticks(classes)  
    plt.grid(axis="y", linestyle="--", alpha=0.6)  
  
    for i, count in enumerate(counts):  
        plt.text(classes[i], count + 5, str(count), ha="center", fontsize=12)  
  
    plt.show()  
  
plot_class_distribution(class_counts)
```



```
In [7]: from sklearn.utils.class_weight import compute_class_weight
import numpy as np

all_labels = []
for _, labels in train_dataset:
    all_labels.extend(np.argmax(labels.numpy(), axis=1))

all_labels = np.array(all_labels)

class_weights_values = compute_class_weight(
    class_weight="balanced",
    classes=np.unique(all_labels),
    y=all_labels
)

class_weights = {i: class_weights_values[i] for i in np.unique(all_labels)}

print("Computed Class Weights:", class_weights)
```

Computed Class Weights: {0: 2.7730272596843615, 1: 1.2031123560535326, 2: 1.4319158393836124, 3: 1.2107999749420535, 4: 1.3848248190871963, 5: 1.4557505460570912, 6: 9.168880455407969, 7: 0.6617818256522632, 8: 1.43372153401083, 9: 0.7470047151580738, 10: 1.4727217311795184, 11: 0.7895424836601307, 12: 0.8317053229484918, 13: 1.026120195370567, 14: 0.26651223077135217, 15: 3.8024788510722014, 16: 0.9250945292681759}

data augmentation

```
In [8]: from tensorflow.keras.layers import RandomFlip, RandomRotation, RandomZoom

data_augmentation = tf.keras.Sequential([
    RandomFlip("horizontal"),
    RandomRotation(0.2),
    RandomZoom(0.2),
])
```

```
In [9]: train_dataset = train_dataset.map(lambda x, y: (data_augmentation(x, training=True), y))
```

vgg19

```
In [13]: from tensorflow.keras.applications import VGG19
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D, Dropout
from tensorflow.keras.regularizers import l2
from tensorflow.keras.optimizers import AdamW
import tensorflow as tf
```

```
In [14]: base_model_v = VGG19(weights="imagenet", include_top=False, input_shape=(224, 224, 3))

base_model_v.trainable = False

x = GlobalAveragePooling2D()(base_model_v.output)
x = Dense(256, activation="relu", kernel_regularizer=l2(0.001))(x)
x = Dropout(0.5)(x)
x = Dense(128, activation="relu", kernel_regularizer=l2(0.001))(x)
x = Dropout(0.3)(x)
output = Dense(17, activation="softmax")(x)

model_v = Model(inputs=base_model_v.input, outputs=output)
```

```
In [17]: model_v.compile(optimizer=AdamW(learning_rate=1e-3),
                        loss="categorical_crossentropy",
                        metrics=["accuracy"])

early_stopping = tf.keras.callbacks.EarlyStopping(monitor="val_loss", patience=5, restore_best_weights=True)
lr_scheduler = tf.keras.callbacks.ReduceLROnPlateau(monitor="val_loss", patience=3, factor=0.5)
```

```
In [12]: history = model_v.fit(train_dataset,
                                validation_data=val_dataset,
                                epochs=30,
                                class_weight=class_weights,
                                callbacks=[early_stopping, lr_scheduler])
```

Epoch 1/30
604/604 ————— **8152s** 13s/step - accuracy: 0.3474 - loss: 2.8529 - val_accuracy: 0.7504 - val_loss: 1.1389 - learning_rate: 0.0010

Epoch 2/30
604/604 ————— **7897s** 13s/step - accuracy: 0.6543 - loss: 1.3760 - val_accuracy: 0.7887 - val_loss: 0.9182 - learning_rate: 0.0010

Epoch 3/30
604/604 ————— **7761s** 13s/step - accuracy: 0.7125 - loss: 1.1573 - val_accuracy: 0.8058 - val_loss: 0.8405 - learning_rate: 0.0010

Epoch 4/30
604/604 ————— **7719s** 13s/step - accuracy: 0.7502 - loss: 1.0203 - val_accuracy: 0.8317 - val_loss: 0.7416 - learning_rate: 0.0010

Epoch 5/30
604/604 ————— **7731s** 13s/step - accuracy: 0.7517 - loss: 0.9646 - val_accuracy: 0.8392 - val_loss: 0.6908 - learning_rate: 0.0010

Epoch 6/30
604/604 ————— **7759s** 13s/step - accuracy: 0.7721 - loss: 0.8976 - val_accuracy: 0.8554 - val_loss: 0.6493 - learning_rate: 0.0010

Epoch 7/30
604/604 ————— **7753s** 13s/step - accuracy: 0.7799 - loss: 0.8616 - val_accuracy: 0.8433 - val_loss: 0.6706 - learning_rate: 0.0010

Epoch 8/30
604/604 ————— **7712s** 13s/step - accuracy: 0.7812 - loss: 0.8573 - val_accuracy: 0.8438 - val_loss: 0.6379 - learning_rate: 0.0010

Epoch 9/30
604/604 ————— **7735s** 13s/step - accuracy: 0.7931 - loss: 0.8310 - val_accuracy: 0.8496 - val_loss: 0.6541 - learning_rate: 0.0010

Epoch 10/30
604/604 ————— **7739s** 13s/step - accuracy: 0.7929 - loss: 0.8222 - val_accuracy: 0.8633 - val_loss: 0.6212 - learning_rate: 0.0010

Epoch 11/30
604/604 ————— **7736s** 13s/step - accuracy: 0.7891 - loss: 0.8209 - val_accuracy: 0.8375 - val_loss: 0.7070 - learning_rate: 0.0010

Epoch 12/30
604/604 ————— **7852s** 13s/step - accuracy: 0.7979 - loss: 0.8047 - val_accuracy: 0.8313 - val_loss: 0.7080 - learning_rate: 0.0010

Epoch 13/30
604/604 ————— **8135s** 13s/step - accuracy: 0.7980 - loss: 0.8122 - val_accuracy: 0.8617 - val_loss: 0.6148 - learning_rate: 0.0010

Epoch 14/30
604/604 ————— **7798s** 13s/step - accuracy: 0.8007 - loss: 0.7800 - val_accuracy: 0.8763 - val_loss: 0.5705 - learning_rate: 0.0010

Epoch 15/30
604/604 ————— **7809s** 13s/step - accuracy: 0.8021 - loss: 0.7916 - val_accuracy: 0.8554 - val_loss: 0.6347 - learning_rate: 0.0010

Epoch 16/30
604/604 ————— **7795s** 13s/step - accuracy: 0.8056 - loss: 0.7940 - val_accuracy: 0.8346 - val_loss: 0.6774 - learning_rate: 0.0010

Epoch 17/30
604/604 ————— **7771s** 13s/step - accuracy: 0.7972 - loss: 0.8014 - val_accuracy: 0.8708 - val_loss: 0.5781 - learning_rate: 0.0010

Epoch 18/30
604/604 ————— **7768s** 13s/step - accuracy: 0.8269 - loss: 0.7143 - val_accuracy: 0.8746 - val_loss: 0.5686 - learning_rate: 5.0000e-04









Epoch 19/30
604/604 ————— **7768s** 13s/step - accuracy: 0.8345 - loss: 0.6792 - val_accuracy: 0.8642 - val_loss: 0.5842 - learning_rate: 5.0000e-04

Epoch 20/30
604/604 ————— **7791s** 13s/step - accuracy: 0.8317 - loss: 0.6709 - val_accuracy: 0.8767 - val_loss: 0.5374 - learning_rate: 5.0000e-04

Epoch 21/30
604/604 ————— **7673s** 13s/step - accuracy: 0.8368 - loss: 0.6505 - val_accuracy: 0.8896 - val_loss: 0.4952 - learning_rate: 5.0000e-04

Epoch 22/30
604/604 ————— **7632s** 13s/step - accuracy: 0.8344 - loss: 0.6499 - val_accuracy: 0.8679 - val_loss: 0.5378 - learning_rate: 5.0000e-04

Epoch 23/30

604/604  **7820s** 13s/step - accuracy: 0.8424 - loss: 0.6395 - val_accuracy: 0.8825 - val_loss: 0.5144 - learning_rate: 5.0000e-04
 Epoch 24/30
604/604  **7809s** 13s/step - accuracy: 0.8407 - loss: 0.6341 - val_accuracy: 0.8792 - val_loss: 0.5245 - learning_rate: 5.0000e-04
 Epoch 25/30
604/604  **7831s** 13s/step - accuracy: 0.8484 - loss: 0.5926 - val_accuracy: 0.8883 - val_loss: 0.4891 - learning_rate: 2.5000e-04
 Epoch 26/30
604/604  **8868s** 15s/step - accuracy: 0.8526 - loss: 0.5703 - val_accuracy: 0.8858 - val_loss: 0.4797 - learning_rate: 2.5000e-04
 Epoch 27/30
604/604  **7691s** 13s/step - accuracy: 0.8593 - loss: 0.5734 - val_accuracy: 0.8850 - val_loss: 0.4910 - learning_rate: 2.5000e-04
 Epoch 28/30
604/604  **7701s** 13s/step - accuracy: 0.8593 - loss: 0.5624 - val_accuracy: 0.8933 - val_loss: 0.4591 - learning_rate: 2.5000e-04
 Epoch 29/30
604/604  **7710s** 13s/step - accuracy: 0.8593 - loss: 0.5537 - val_accuracy: 0.8967 - val_loss: 0.4548 - learning_rate: 2.5000e-04
 Epoch 30/30
604/604  **7757s** 13s/step - accuracy: 0.8569 - loss: 0.5565 - val_accuracy: 0.8988 - val_loss: 0.4529 - learning_rate: 2.5000e-04

```

In [21]: # Unfreeze Last 80 Layers for Fine-Tuning
         for layer in base_model_v.layers[-80:]:
             layer.trainable = True

         # Compile again with a lower learning rate
         model_v.compile(optimizer=AdamW(learning_rate=1e-5),
                        loss="categorical_crossentropy",
                        metrics=["accuracy"])
  
```



```
In [22]: fine_tune_history = model_v.fit(train_dataset,
                                         validation_data=val_dataset,
                                         epochs=30,
                                         class_weight=class_weights,
                                         callbacks=[early_stopping, lr_scheduler])
```

Epoch 1/30
604/604 ————— **37441s** 62s/step - accuracy: 0.8809 - loss: 0.5124 - val_accuracy: 0.9133 - val_loss: 0.3677 - learning_rate: 1.0000e-05

Epoch 2/30
604/604 ————— **37324s** 62s/step - accuracy: 0.9098 - loss: 0.4103 - val_accuracy: 0.9429 - val_loss: 0.3114 - learning_rate: 1.0000e-05

Epoch 3/30
604/604 ————— **37414s** 62s/step - accuracy: 0.9303 - loss: 0.3437 - val_accuracy: 0.9371 - val_loss: 0.3127 - learning_rate: 1.0000e-05

Epoch 4/30
604/604 ————— **37556s** 62s/step - accuracy: 0.9411 - loss: 0.3054 - val_accuracy: 0.9625 - val_loss: 0.2394 - learning_rate: 1.0000e-05

Epoch 5/30
604/604 ————— **36889s** 61s/step - accuracy: 0.9458 - loss: 0.2999 - val_accuracy: 0.9442 - val_loss: 0.2923 - learning_rate: 1.0000e-05

Epoch 6/30
604/604 ————— **36366s** 60s/step - accuracy: 0.9531 - loss: 0.2786 - val_accuracy: 0.9400 - val_loss: 0.3015 - learning_rate: 1.0000e-05

Epoch 7/30
604/604 ————— **36180s** 60s/step - accuracy: 0.9576 - loss: 0.2709 - val_accuracy: 0.9613 - val_loss: 0.2429 - learning_rate: 1.0000e-05

Epoch 8/30
604/604 ————— **36988s** 61s/step - accuracy: 0.9660 - loss: 0.2348 - val_accuracy: 0.9683 - val_loss: 0.2249 - learning_rate: 5.0000e-06

Epoch 9/30
604/604 ————— **37462s** 62s/step - accuracy: 0.9716 - loss: 0.2165 - val_accuracy: 0.9667 - val_loss: 0.2261 - learning_rate: 5.0000e-06

Epoch 10/30
604/604 ————— **37451s** 62s/step - accuracy: 0.9712 - loss: 0.2140 - val_accuracy: 0.9588 - val_loss: 0.2621 - learning_rate: 5.0000e-06

Epoch 11/30
604/604 ————— **37444s** 62s/step - accuracy: 0.9719 - loss: 0.2195 - val_accuracy: 0.9646 - val_loss: 0.2421 - learning_rate: 5.0000e-06

Epoch 12/30
604/604 ————— **37423s** 62s/step - accuracy: 0.9808 - loss: 0.1861 - val_accuracy: 0.9563 - val_loss: 0.2818 - learning_rate: 2.5000e-06

Epoch 13/30
604/604 ————— **37414s** 62s/step - accuracy: 0.9812 - loss: 0.1901 - val_accuracy: 0.9717 - val_loss: 0.2155 - learning_rate: 2.5000e-06

Epoch 14/30
604/604 ————— **37425s** 62s/step - accuracy: 0.9828 - loss: 0.1756 - val_accuracy: 0.9675 - val_loss: 0.2487 - learning_rate: 2.5000e-06

Epoch 15/30
604/604 ————— **37422s** 62s/step - accuracy: 0.9815 - loss: 0.1773 - val_accuracy: 0.9667 - val_loss: 0.2384 - learning_rate: 2.5000e-06

Epoch 16/30
604/604 ————— **37423s** 62s/step - accuracy: 0.9817 - loss: 0.1797 - val_accuracy: 0.9700 - val_loss: 0.2195 - learning_rate: 2.5000e-06

Epoch 17/30
604/604 ————— **37407s** 62s/step - accuracy: 0.9847 - loss: 0.1697 - val_accuracy: 0.9754 - val_loss: 0.2090 - learning_rate: 1.2500e-06

Epoch 18/30
604/604 ————— **37431s** 62s/step - accuracy: 0.9887 - loss: 0.1591 - val_accuracy: 0.9842 - val_loss: 0.1801 - learning_rate: 1.2500e-06


Epoch 19/30
604/604 ————— **41204s** 68s/step - accuracy: 0.9873 - loss: 0.1653 - val_accuracy: 0.9787 - val_loss: 0.1946 - learning_rate: 1.2500e-06

Epoch 20/30
604/604 ————— **37535s** 62s/step - accuracy: 0.9877 - loss: 0.1631 - val_accuracy: 0.9771 - val_loss: 0.2084 - learning_rate: 1.2500e-06

Epoch 21/30
604/604 ————— **37473s** 62s/step - accuracy: 0.9874 - loss: 0.1638 - val_accuracy: 0.9729 - val_loss: 0.2279 - learning_rate: 1.2500e-06

Epoch 22/30
604/604 ————— **37429s** 62s/step - accuracy: 0.9875 - loss: 0.1630 - val_accuracy: 0.9750 - val_loss: 0.2127 - learning_rate: 1.0000e-06

Epoch 23/30

604/604  **37456s** 62s/step - accuracy: 0.9890 - loss: 0.1561 - val_accuracy: 0.9775 - val_loss: 0.2082 - learning_rate: 1.0000e-06

In [15]: `model_v.save("vgg19.h5")`

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

In [17]: `model_v.save('my_model.keras')`


```
In [20]: from tensorflow.keras.models import load_model

# Load the model
model_v = load_model("my_model.keras")

# Check model summary
model_v.summary()
```

Model: "functional_1"

Layer (type)	Output Shape	Param #
input_layer_1 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1,792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36,928
block1_pool1 (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73,856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147,584
block2_pool1 (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295,168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590,080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590,080
block3_conv4 (Conv2D)	(None, 56, 56, 256)	590,080
block3_pool1 (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1,180,160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_conv4 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_pool1 (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv4 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_pool1 (MaxPooling2D)	(None, 7, 7, 512)	0
global_average_pooling2d (GlobalAveragePooling2D)	(None, 512)	0
dense (Dense)	(None, 256)	131,328
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 128)	32,896
dropout_1 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 17)	2,193



Total params: 40,381,604 (154.04 MB)

Trainable params: 20,190,801 (77.02 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 20,190,803 (77.02 MB)

```
In [24]: model_v.save("vgg19.keras")
```

```
In [25]: test_loss, test_acc = model_v.evaluate(test_dataset)
print(f"Test Accuracy: {test_acc:.4f}")
```

77/77 ————— 909s 11s/step - accuracy: 0.9769 - loss: 0.1941
Test Accuracy: 0.9811

```
In [26]: train_loss, train_acc = model_v.evaluate(train_dataset)
print(f"Train Accuracy: {train_acc:.4f}")
```

604/604 ————— 6955s 12s/step - accuracy: 0.9954 - loss: 0.1367
Train Accuracy: 0.9960

```
In [27]: val_loss, val_acc = model_v.evaluate(val_dataset)
print(f"Validation Accuracy: {val_acc:.4f}")
```

75/75 ————— 878s 11s/step - accuracy: 0.9849 - loss: 0.1799
Validation Accuracy: 0.9821


```
In [28]: def plot_metrics(history):
plt.figure(figsize=(12, 4))

# Accuracy Plot
plt.subplot(1, 2, 1)
plt.plot(history.history["accuracy"], label="Train Accuracy")
plt.plot(history.history["val_accuracy"], label="Validation Accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.title("Training vs Validation Accuracy")
plt.legend()

# Loss Plot
plt.subplot(1, 2, 2)
plt.plot(history.history["loss"], label="Train Loss")
plt.plot(history.history["val_loss"], label="Validation Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Training vs Validation Loss")
plt.legend()

plt.show()

plot_metrics(fine_tune_history)
```



```
In [29]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix, classification_report

true_labels = []
pred_labels = []

for images, labels in test_dataset:
    preds = model_v.predict(images)
    pred_classes = np.argmax(preds, axis=1)
    true_classes = np.argmax(labels.numpy(), axis=1)

    true_labels.extend(true_classes)
    pred_labels.extend(pred_classes)

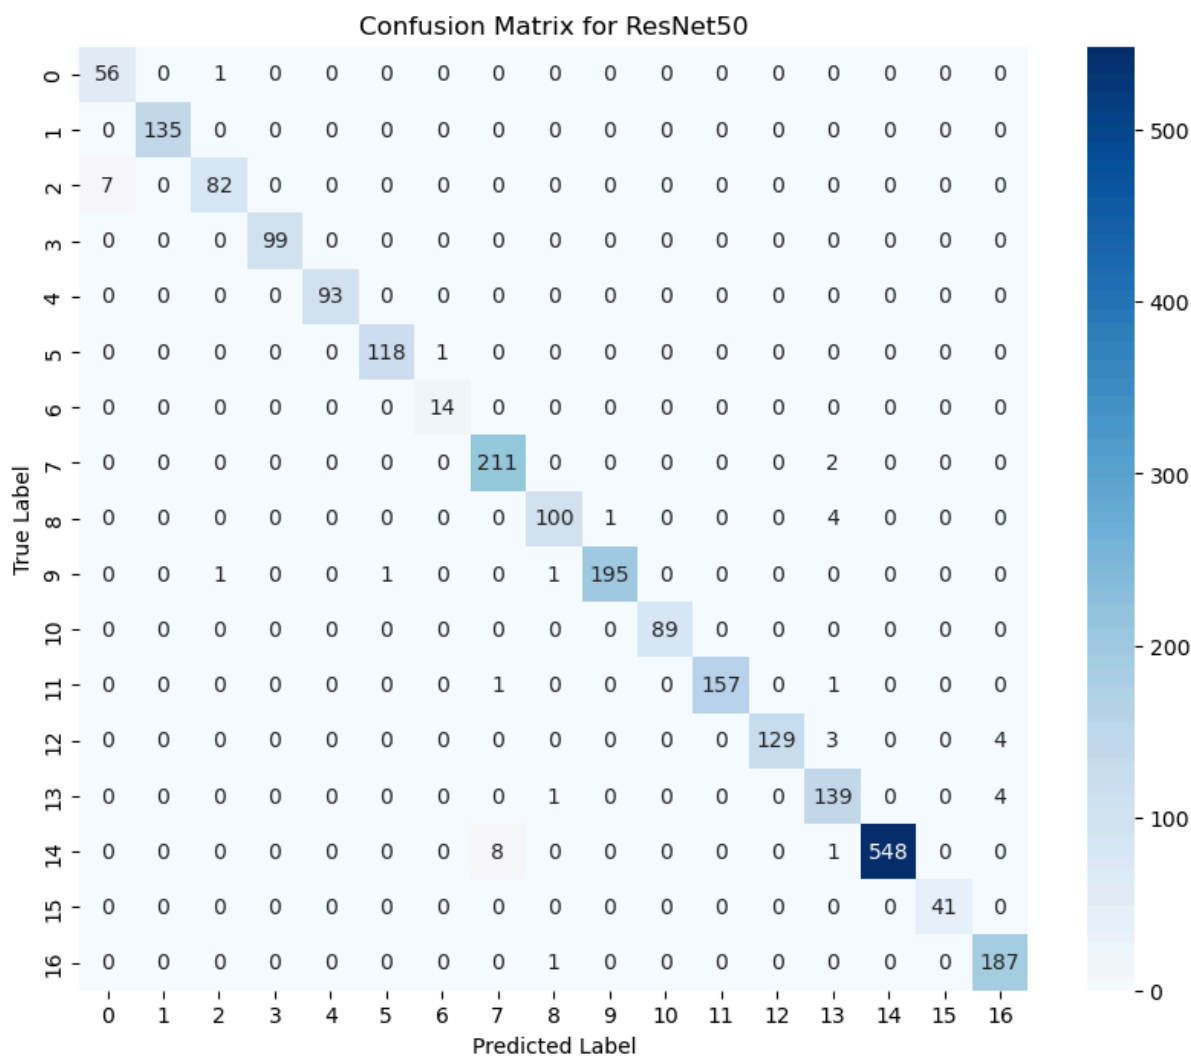
true_labels = np.array(true_labels)
pred_labels = np.array(pred_labels)

cm = confusion_matrix(true_labels, pred_labels)
```


[illegible]

```
1/1 ----- 11s 11s/step
1/1 ----- 11s 11s/step
1/1 ----- 12s 12s/step
1/1 ----- 11s 11s/step
1/1 ----- 11s 11s/step
1/1 ----- 11s 11s/step
1/1 ----- 11s 11s/step
1/1 ----- 11s 11s/step
1/1 ----- 11s 11s/step
1/1 ----- 11s 11s/step
1/1 ----- 2s 2s/step
```

```
In [30]: plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=range(17), yticklabels=range(17))
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix for ResNet50")
plt.show()
```



```
In [31]: #classification report
class_names = [f"Class {i}" for i in range(17)]
report = classification_report(true_labels, pred_labels, target_names=class_names)
print("Classification Report:\n", report)
```

Classification Report:

	precision	recall	f1-score	support
Class 0	0.89	0.98	0.93	57
Class 1	1.00	1.00	1.00	135
Class 2	0.98	0.92	0.95	89
Class 3	1.00	1.00	1.00	99
Class 4	1.00	1.00	1.00	93
Class 5	0.99	0.99	0.99	119
Class 6	0.93	1.00	0.97	14
Class 7	0.96	0.99	0.97	213
Class 8	0.97	0.95	0.96	105
Class 9	0.99	0.98	0.99	198
Class 10	1.00	1.00	1.00	89
Class 11	1.00	0.99	0.99	159
Class 12	1.00	0.95	0.97	136
Class 13	0.93	0.97	0.95	144
Class 14	1.00	0.98	0.99	557
Class 15	1.00	1.00	1.00	41
Class 16	0.96	0.99	0.98	188
accuracy			0.98	2436
macro avg	0.98	0.98	0.98	2436
weighted avg	0.98	0.98	0.98	2436

In []: